

## Suman Acharya

Q1) Use the CoNLL-2003 Named Entity Recognition dataset which contains four entity types: PER (Person names) LOC (Locations) ORG (Organizations) MISC (Miscellaneous entities)  
Load the CoNLL-2003 dataset using HuggingFace datasets (<https://huggingface.co/datasets/eriktks/conll2003>) and initialize a Weights & Biases project called "Q1-weak-supervision-ner". Log the dataset statistics (number of samples, entity distribution) to W&B as summary metrics

```
(Assignment5_mlops) PS D:\Assignment5_mlops> python Q1.py
/ag4kog1d
Number of training samples: 14987
Number of validation samples: 3684
Number of test samples: 3466
Entity distribution: {'O': 170524, 'B-ORG': 6321, 'B-MISC': 3438, 'B-PER': 6600, 'I-PER': 4528, 'B-LOC': 7140, 'I-ORG': 3704, 'I-MISC': 1155, 'I-LOC': 1157}
wandb:
wandb: Run summary:
wandb:   num_test_samples 3466
wandb:   num_train_samples 14987
wandb: num_validation_samples 3684
wandb:
wandb: View run conll2003-local-stats at: https://wandb.ai/142402011-indian-institute-of-technology/Q1-weak-supervision-ner/runs/ag4kog1d
wandb: View project at: https://wandb.ai/142402011-indian-institute-of-technology/Q1-weak-supervision-ner
wandb: Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)
wandb: Find logs at: .\wandb\run-20251013_125300-ag4kog1d\logs
(Assignment5_mlops) PS D:\Assignment5_mlops>
```

## Summary

Summary metrics are your model's outputs. [Learn more](#)

🔍 Search keys with regex

▼ **Summary metrics:** {} 12 keys

entity\_distribution.B-LOC: 7,140

entity\_distribution.B-MISC: 3,438

entity\_distribution.B-ORG: 6,321

entity\_distribution.B-PER: 6,600

entity\_distribution.I-LOC: 1,157

entity\_distribution.I-MISC: 1,155

entity\_distribution.I-ORG: 3,704

entity\_distribution.I-PER: 4,528

entity\_distribution.O: 170,524

num\_test\_samples: 3,466

num\_train\_samples: 14,987

num\_validation\_samples: 3,684

Q2) Implement two basic labeling functions using Snorkel AI: a. A heuristic function detecting years (1900-2099) as potential DATE/MISC entities b. A pattern-matching function identifying organizations by common suffixes ("Inc.", "Corp.", "Ltd.") Log each labeling function's coverage and accuracy to W&B using wandb.log()

```

OUTPUT  DEBUG CONSOLE  TERMINAL  PORTS
(Assignment5_mlops) PS D:\Assignment5_mlops> python Q2.py
wandb: Syncing run conll2003-lf-analysis-int
wandb: View project at https://wandb.ai/142402011-indian-institute-of-technology/Q1-weak-supervision-ner
wandb: View run at https://wandb.ai/142402011-indian-institute-of-technology/Q1-weak-supervision-ner/runs/30d70a301008 | 14987/14987 [00:01<00
Labeling Function Summary:
      j Polarity Coverage Overlaps Conflicts
lf_year      0      [0] 0.095549 0.000133 0.000133
lf_org_suffix 1      [1] 0.001201 0.000133 0.000133
wandb:
wandb: Run history:
wandb: lf_org_suffix_coverage _
wandb: lf_year_coverage _
wandb:
wandb: Run summary:
wandb: lf_org_suffix_coverage 0.0012
wandb: lf_year_coverage 0.09555
wandb:

```

### Q3) Implement Snorkell's Label aggregation (Majority Label Voter)

```

00 san counts = df.tosql()["requested_label"] using counts()
OUTPUT DEBUG CONSOLE TERMINAL PORTS powershell + - [X] [X] [X] [X]
(Assignment5_mlops) PS D:\Assignment5_mlops> python Q3.py
100%|
Labeling Function Summary:
      j Polarity Coverage Overlaps Conflicts
lf_year      0      [0]  0.095549  0.000133  0.000133
lf_org_suffix 1      [1]  0.001201  0.000133  0.000133

Aggregated Label Distribution:
aggregated_label
-1      13541
0        1430
1         16
Name: count, dtype: int64
wandb:
wandb: Run history:
wandb: lf_org_suffix_coverage _
wandb: lf_year_coverage _
wandb:
wandb: Run summary:
wandb: lf_org_suffix_coverage 0.0012
wandb: lf_year_coverage 0.09555
wandb:
wandb: View run conll2003-lf-aggregation at: https://wandb.ai/142402011-indian-institute-of-technology/Q1-weak-supervision-ner/runs/sk32gd8i
wandb: View project at: https://wandb.ai/142402011-indian-institute-of-technology/Q1-weak-supervision-ner
wandb: Synced 5 W&B file(s), 0 media file(s), 0 artifact file(s) and 0 other file(s)
wandb: Find logs at: .\wandb\run-20251013_145703-sk32gd8i\logs
(Assignment5_mlops) PS D:\Assignment5_mlops>

```

View raw data

Summary

View raw data

Summary metrics are your model's outputs. [Learn more](#)

▼ **Summary metrics:** {} 5 keys

aggregated\_label\_distribution.-1: 13,541

aggregated\_label\_distribution.0: 1,430

aggregated\_label\_distribution.1: 16

lf\_org\_suffix\_coverage: 0.0012010409021151664

lf\_year\_coverage: 0.09554947621271768

- Q4) Implement the following in Weights and Bias:
- Train CIFAR 100 and CIFAR 10 sequentially for 100 epochs

```

OUTPUT  DEBUG CONSOLE  TERMINAL  PORTS
(Assignment5_mlops) PS D:\Assignment5_mlops> python Q4.py

Training on CIFAR-100...
[CIFAR100] Epoch 1/20 | Loss: 4.368 | Acc: 5.46%
[CIFAR100] Epoch 2/20 | Loss: 3.800 | Acc: 11.20%
[CIFAR100] Epoch 3/20 | Loss: 3.481 | Acc: 16.42%
[CIFAR100] Epoch 4/20 | Loss: 3.172 | Acc: 20.96%
[CIFAR100] Epoch 5/20 | Loss: 2.832 | Acc: 27.42%
[CIFAR100] Epoch 6/20 | Loss: 2.449 | Acc: 34.46%
[CIFAR100] Epoch 7/20 | Loss: 2.090 | Acc: 42.98%
[CIFAR100] Epoch 8/20 | Loss: 1.763 | Acc: 50.30%
[CIFAR100] Epoch 9/20 | Loss: 1.240 | Acc: 64.40%
[CIFAR100] Epoch 10/20 | Loss: 1.013 | Acc: 70.48%
[CIFAR100] Epoch 11/20 | Loss: 0.867 | Acc: 74.98%
[CIFAR100] Epoch 12/20 | Loss: 0.623 | Acc: 81.78%
[CIFAR100] Epoch 13/20 | Loss: 0.343 | Acc: 90.56%
[CIFAR100] Epoch 14/20 | Loss: 0.223 | Acc: 93.92%
[CIFAR100] Epoch 15/20 | Loss: 0.262 | Acc: 92.56%
[CIFAR100] Epoch 16/20 | Loss: 0.255 | Acc: 92.90%
[CIFAR100] Epoch 17/20 | Loss: 0.296 | Acc: 92.04%
[CIFAR100] Epoch 18/20 | Loss: 0.327 | Acc: 90.58%
[CIFAR100] Epoch 19/20 | Loss: 0.365 | Acc: 89.48%
[CIFAR100] Epoch 20/20 | Loss: 0.154 | Acc: 95.64%
  
```

```
(Assignment5_mlops) PS D:\Assignment5_mlops> python Q4.py
```

```
Fine-tuning on CIFAR-10...
```

[CIFAR10]	Epoch 1/20	Loss: 1.652	Acc: 41.90%
[CIFAR10]	Epoch 2/20	Loss: 1.094	Acc: 60.52%
[CIFAR10]	Epoch 3/20	Loss: 0.755	Acc: 73.26%
[CIFAR10]	Epoch 4/20	Loss: 0.572	Acc: 80.54%
[CIFAR10]	Epoch 5/20	Loss: 0.404	Acc: 85.94%
[CIFAR10]	Epoch 6/20	Loss: 0.295	Acc: 89.76%
[CIFAR10]	Epoch 7/20	Loss: 0.190	Acc: 93.40%
[CIFAR10]	Epoch 8/20	Loss: 0.130	Acc: 96.04%
[CIFAR10]	Epoch 9/20	Loss: 0.340	Acc: 89.72%
[CIFAR10]	Epoch 10/20	Loss: 0.264	Acc: 91.16%
[CIFAR10]	Epoch 11/20	Loss: 0.106	Acc: 96.66%
[CIFAR10]	Epoch 12/20	Loss: 0.136	Acc: 95.82%
[CIFAR10]	Epoch 13/20	Loss: 0.103	Acc: 96.66%
[CIFAR10]	Epoch 14/20	Loss: 0.068	Acc: 97.84%
[CIFAR10]	Epoch 15/20	Loss: 0.048	Acc: 98.46%
[CIFAR10]	Epoch 16/20	Loss: 0.041	Acc: 98.94%
[CIFAR10]	Epoch 17/20	Loss: 0.104	Acc: 96.62%
[CIFAR10]	Epoch 18/20	Loss: 0.207	Acc: 93.40%
[CIFAR10]	Epoch 19/20	Loss: 0.141	Acc: 95.46%
[CIFAR10]	Epoch 20/20	Loss: 0.089	Acc: 97.20%

```
wandb: Run history:
```

```
wandb: CIFAR100_accuracy
```

```
wandb: CIFAR100_loss
```

```
wandb: CIFAR10_accuracy
```

```
wandb: CIFAR10_loss
```

```
wandb: epoch
```

```
wandb:
```

```
wandb: Run summary:
```

```
wandb: CIFAR100_accuracy 95.64
```

```
wandb: CIFAR100_loss 0.15435
```

```
wandb: CIFAR10_accuracy 97.2
```

```
wandb: CIFAR10_loss 0.08921
```

```
wandb: epoch 20
```

## Config

[View raw data](#)Config parameters are your model's inputs. [Learn more](#) Search keys with regex

## ▼ Config parameters: {} 4 keys

batch\_size: 64

epochs: 20

lr: 0.001

subset\_ratio: 0.1

## Summary

[View raw data](#)Summary metrics are your model's outputs. [Learn more](#) Search keys with regex

## ▼ Summary metrics: {} 5 keys

CIFAR10\_accuracy: 97.2

CIFAR10\_loss: 0.0892089693748121

CIFAR100\_accuracy: 95.64

CIFAR100\_loss: 0.1543478622466703

epoch: 20

b. Train CIFAR 10 and CIFAR 100 sequentially for 100 epochs.

(Assignment5\_mlops) PS D:\Assignment5\_mlops&gt; python Q4.py

Training on CIFAR-10...

[CIFAR10]	Epoch 1/20	Loss: 1.884	Acc: 33.38%
[CIFAR10]	Epoch 2/20	Loss: 1.507	Acc: 45.60%
[CIFAR10]	Epoch 3/20	Loss: 1.326	Acc: 52.90%
[CIFAR10]	Epoch 4/20	Loss: 1.128	Acc: 59.62%
[CIFAR10]	Epoch 5/20	Loss: 0.986	Acc: 65.66%
[CIFAR10]	Epoch 6/20	Loss: 0.816	Acc: 70.60%
[CIFAR10]	Epoch 7/20	Loss: 0.667	Acc: 76.00%
[CIFAR10]	Epoch 8/20	Loss: 0.621	Acc: 79.08%
[CIFAR10]	Epoch 9/20	Loss: 0.544	Acc: 81.02%
[CIFAR10]	Epoch 10/20	Loss: 0.362	Acc: 87.86%
[CIFAR10]	Epoch 11/20	Loss: 0.304	Acc: 89.22%
[CIFAR10]	Epoch 12/20	Loss: 0.265	Acc: 91.20%
[CIFAR10]	Epoch 13/20	Loss: 0.319	Acc: 89.22%
[CIFAR10]	Epoch 14/20	Loss: 0.213	Acc: 92.64%
[CIFAR10]	Epoch 15/20	Loss: 0.168	Acc: 95.20%
[CIFAR10]	Epoch 16/20	Loss: 0.403	Acc: 86.74%
[CIFAR10]	Epoch 17/20	Loss: 0.094	Acc: 96.72%
[CIFAR10]	Epoch 18/20	Loss: 0.057	Acc: 98.62%
[CIFAR10]	Epoch 19/20	Loss: 0.140	Acc: 95.20%
[CIFAR10]	Epoch 20/20	Loss: 0.110	Acc: 96.30%

Fine-tuning on CIFAR-100...

[CIFAR100]	Epoch 1/20	Loss: 4.128	Acc: 7.68%
[CIFAR100]	Epoch 2/20	Loss: 3.299	Acc: 20.12%
[CIFAR100]	Epoch 3/20	Loss: 2.697	Acc: 30.20%
[CIFAR100]	Epoch 4/20	Loss: 2.081	Acc: 43.58%
[CIFAR100]	Epoch 5/20	Loss: 1.519	Acc: 57.68%
[CIFAR100]	Epoch 6/20	Loss: 1.046	Acc: 69.20%
[CIFAR100]	Epoch 7/20	Loss: 0.629	Acc: 81.58%
[CIFAR100]	Epoch 8/20	Loss: 0.395	Acc: 89.08%
[CIFAR100]	Epoch 9/20	Loss: 0.325	Acc: 90.94%
[CIFAR100]	Epoch 10/20	Loss: 0.227	Acc: 93.64%
[CIFAR100]	Epoch 11/20	Loss: 0.368	Acc: 89.24%
[CIFAR100]	Epoch 12/20	Loss: 0.307	Acc: 90.84%
[CIFAR100]	Epoch 13/20	Loss: 0.158	Acc: 96.10%
[CIFAR100]	Epoch 14/20	Loss: 0.176	Acc: 94.86%
[CIFAR100]	Epoch 15/20	Loss: 0.115	Acc: 97.10%
[CIFAR100]	Epoch 16/20	Loss: 0.145	Acc: 96.16%
[CIFAR100]	Epoch 17/20	Loss: 0.135	Acc: 96.38%
[CIFAR100]	Epoch 18/20	Loss: 0.159	Acc: 95.98%
[CIFAR100]	Epoch 19/20	Loss: 0.258	Acc: 92.20%
[CIFAR100]	Epoch 20/20	Loss: 0.136	Acc: 95.90%

andb: Run history:

andb: CIFAR100\_accuracy

andb: CIFAR100\_loss

andb: CIFAR10\_accuracy

andb: CIFAR10\_loss

andb: epoch

andb:

andb: Run summary:

andb: CIFAR100\_accuracy 95.9

andb: CIFAR100\_loss 0.13563

andb: CIFAR10\_accuracy 96.3

andb: CIFAR10\_loss 0.11019

andb: epoch 20

## Config

[View raw data](#)Config parameters are your model's inputs. [Learn more](#)

No config parameters were saved for this run.

Check the [configuration documentation](#) for more information.

## Summary

Summary metrics are your model's outputs. [Learn more](#) Search keys with regex

## ▼ Summary metrics: {} 5 keys

CIFAR10\_accuracy: 96.3

CIFAR10\_loss: 0.1101879781654364

CIFAR100\_accuracy: 95.9

CIFAR100\_loss: 0.1356346655071159

epoch: 20

**Write your observations with experimental proof.**

From the above two result we can see that when we train our model with CIFAR100 and then fine tuning it with CIFAR10 it gives better result and also faster convergence compare to we train with CIFAR10 and the fine tuning with CIFAR100 .

**Reason**

CIFAR-100 has more classes and complex .features; pretraining learns diverse representations that help CIFAR-10

CIFAR-10 features are simpler; transferring to CIFAR-100 (more complex task) is harder